

PXYZ System Reference: Comprehensive Guide

Executive Summary

PXYZ is a revolutionary **workflow compiler and runtime system** designed to safely grant AI agents access to production systems. Unlike traditional imperative programming, PXYZ treats business logic as **auditable data** rather than executable code.

Core Value Proposition:

- **Provable Safety:** Workflows are finite, deterministic, and bounded
- **Full Auditability:** Every operation is explicit and reviewable
- **Compile-Time Safety:** Dangerous patterns are caught before deployment
- **Human-in-the-Loop:** Irreversible actions require human confirmation
- **Production-Ready:** ~600 lines of verifiable WebAssembly runtime

Part 1: System Architecture

Three-Component Design

1. Rust Compiler (~1,500 LOC)

Purpose: Transform XML workflow definitions into portable binary artifacts

Responsibilities:

- Parse XML into Abstract Syntax Tree (AST)
- Lower AST to Graph Intermediate Representation (IR)
- Compile predicates to Predicate VM bytecode
- Apply three-layer validation (Syntactic, Semantic, Pragmatic)
- Optimize for binary size and execution efficiency
- Emit portable `graph.bin` binary

Output: `graph.bin` - A self-contained, portable workflow binary

2. WASM Runtime (~600 LOC)

Purpose: Execute workflows in a completely sandboxed environment

Architecture:

- Written in WebAssembly Text (.wat) for maximum auditability

- Embeds Predicate VM for safe condition evaluation
- Enforces strict execution limits (1000 visited nodes, 256 predicate steps max)
- Zero direct system access (fully sandboxed)
- All side effects mediated through host imports

Safety Limits:

- MAX_VISITED_NODES = 1000 (prevent runaway traversal)
- MAX_PREDICATE_STEPS = 256 (prevent infinite loops)
- MAX_STACK_DEPTH = 16 (prevent stack overflow)
- MAX_CALL_DEPTH = 4 (prevent recursion)
- MAX_PREDICATE_BYTECODE = 256 bytes (limit complexity)

3. JavaScript Host

Purpose: Provide the bridge between sandboxed runtime and external systems

Responsibilities:

- Implement host imports (io_call, io_resolve_var, etc.)
- Resolve variables from application context
- Execute external I/O operations
- Log and audit all operations
- Return responses to runtime

Host Imports:

- io_call(op_code) - Execute external operation
- io_resolve_var(path) - Resolve \$variable.paths
- io_is_human() - Check if actor is human
- io_is_confirmed(entity) - Check confirmation status
- io_log(message) - Emit log messages
- emit_event(event) - Emit audit events

Part 2: Workflow Definition Language (XML)

Document Structure

```
&lt;omar&gt;
  &lt;schemas&gt;&lt;/schemas&gt;
  &lt;predicates&gt;&lt;/predicates&gt;
  &lt;workflow&gt;&lt;/workflow&gt;
```

```
&lt;templates&ampgt&lt;/templates&ampgt  
&lt;/omar&gt;
```

Core Elements

Nodes

Definition: Units of work or control points in the workflow graph

Kind	Purpose	Characteristics
Transform	Validate/transform data	No side effects, local only
External	Call I/O handler	Side effects, irreversible
Render	Generate output	Template-based, user-facing
Signal	UI framework signal	Client-side changes
Auth	Authorization check	Permission validation
Terminal	Success endpoint	Return 2xx status
Error	Error endpoint	Return error status

Edges

Definition: Directed connections between nodes

Properties:

- **from** (required): Source node ID
- **to** (required): Target node ID
- **weight** (optional): Priority (higher = evaluated first)
- **parallel** (optional): Can be traversed in parallel
- **fallback** (optional): Only taken if other paths fail

Conditional Traversal:

```
&lt;edge from="A" to="B"&gt;  
  &lt;when&gt;  
    &lt;and&gt;  
      &lt;eq left="$token.role" right="admin"/&gt;  
      &lt;contains left="$state.users.ids" right="$token.sub"/&gt;  
    &lt;/and&gt;  
  &lt;/when&gt;  
&lt;/edge&gt;
```

Predicates

Definition: Reusable boolean conditions

Available Operations:

Category	Operations
Comparison	eq, neq, gt, gte, lt, lte
String	contains, matches, startsWith, endsWith
Logic	and, or, not
Type	fn name="is_defined", fn name="is_null"

Variable Paths:

- \$token.sub - User ID
- \$token.perms - Permissions array
- \$token.tenant - Tenant ID
- \$entity.owner_id - Entity owner
- \$entity.status - Entity status
- \$input.query - User input
- \$state.node_id.field - Output from node

Part 3: Compilation Pipeline

Six-Stage Transformation

Stage 1: Parsing

- **Input:** XML text
- **Output:** Abstract Syntax Tree (AST)
- **Function:** XML parsing and basic structure validation

Stage 2: Lowering

- **Input:** AST
- **Output:** Graph Intermediate Representation (IR)
- **Function:** Convert high-level AST to graph-centric IR, resolve symbolic names to numeric IDs

Stage 3: Predicate Compilation

- **Input:** Predicate expressions
- **Output:** Predicate VM bytecode
- **Function:** Compile all predicates to efficient bytecode

Stage 4: Validation

- **Input:** Graph IR
- **Output:** Validation diagnostics
- **Function:** Three-layer constraint checking

Stage 5: Optimization

- **Input:** Validated IR
- **Output:** Optimized IR
- **Function:** Dead code elimination, predicate deduplication

Stage 6: Emission

- **Input:** Optimized IR
- **Output:** `graph.bin` binary
- **Function:** Serialize to binary format with header and data sections

Three-Layer Validation System

Syntactic (SYN) - Structure Validation

Ensures graph structure is well-formed:

- SYN001: Edge targets exist
- SYN002: Entry points reference existing nodes
- SYN003: Predicate references exist
- SYN004: No duplicate node IDs
- SYN005: At least one entry point defined
- SYN006: No duplicate entry points
- SYN007: Edge sources exist

Semantic (SEM) - Logic Validation

Ensures graph logic is coherent:

- SEM001: Auth nodes have predicates
- SEM002: External nodes have op codes
- SEM003: Terminal nodes shouldn't have outgoing edges
- SEM004: No cycles (must be DAG)
- SEM005: All nodes reachable from entry
- SEM006: Error nodes have incoming edges
- SEM007: Render nodes have templates

Pragmatic (PRAG) - Business & Safety Rules

Enforces high-level safety policies:

- PRAG001: LLM → Irreversible action paths require validation gate
- PRAG002: Write operations should have error-handling branches
- PRAG003: Irreversible actions require human-in-the-loop
- PRAG004: Irreversible actions require confirmed inputs
- PRAG005: Quarantined data cannot escape to I/O operations

Part 4: Binary Format (graph.bin)

File Structure

Header (96 bytes)

Offset	Size	Field	Description
0x00	4	Magic	0x504E5958 (ASCII: "PXYZ")
0x04	2	Ver.Major	Binary format version
0x06	2	Ver.Minor	
0x08	4	Node count	Total nodes in graph
0x0C	4	Edge count	Total edges in graph
0x10	4	Pred count	Total predicates
0x14	4	String pool size	String pool bytes
0x18	4	Entry count	Number of entry points
0x20	32	Source hash	SHA-256 of source XML
0x40	4	Nodes offset	Byte offset to nodes section
0x44	4	Edges offset	Byte offset to edges section
0x48	4	Predicates offset	Byte offset to predicates
0x4C	4	Strings offset	Byte offset to string pool
0x50	4	Entries offset	Byte offset to entry points

Node Entry (16 bytes)

Offset	Size	Field	Description
0x00	4	Node ID	Index in node array
0x04	1	Kind	0-6 (Transform, External, etc.)
0x05	1	Flags	Bitfield of properties
0x06	2	Op code	I/O operation code
0x08	4	Data offset	String pool offset
0x0C	2	Edge start	Index in edges array
0x0E	2	Edge count	Number of outgoing edges

Edge Entry (12 bytes)

Offset	Size	Field	Description
0x00	4	Target node ID	Destination node
0x04	2	Predicate ID	Condition ID (0=unconditional)
0x06	2	Reserved	Future use
0x08	2	Weight	Priority (higher first)
0x0A	2	Flags	Edge properties

Entry Point Entry (8 bytes)

Offset	Size	Field	Description
0x00	4	PX hash	FNV-1a hash of (P, X)
0x04	4	Node ID	Starting node for this entry

Data Sections

- **Nodes:** Contiguous array of Node Entries (16 bytes each)
- **Edges:** Sorted by source node ID
- **Predicates:** Variable-length bytecode chunks
- **Strings:** Null-terminated UTF-8 strings (deduplicated pool)
- **Entry Points:** Lookup table for (P, X) → Node ID

Part 5: Predicate VM

Bytecode Instruction Set

The Predicate VM is a stack-based interpreter with 26 opcodes:

Stack Operations

- 0x01 PUSH_INT - Push 32-bit integer
- 0x02 PUSH_STR - Push string reference
- 0x03 LOAD_VAR - Load variable from host
- 0x04 LOAD_FIELD - Get field from object

Comparisons

- 0x10 EQ - Equal (pop a, b → push a==b)
- 0x11 NEQ - Not equal
- 0x12 GT - Greater than
- 0x13 GTE - Greater or equal
- 0x14 LT - Less than
- 0x15 LTE - Less or equal

Logical Operations

- 0x20 AND - Logical AND
- 0x21 OR - Logical OR
- 0x22 NOT - Logical NOT

String Operations

- 0x30 CONTAINS - String contains
- 0x31 MATCHES - Regex match
- 0x32 STARTS_WITH - String prefix
- 0x33 ENDS_WITH - String suffix

Type Operations

- 0x40 LEN - String/array length
- 0x41 GET - Array element access
- 0x42 IS_NULL - Check null
- 0x43 IS_DEFINED - Check defined
- 0x44 IS_CONFIRMED - Check confirmation status

Advanced

- 0xF0 CALL_PRED - Call another predicate (by 16-bit ID)
- 0xFF RET - Return boolean value

Part 6: IO Operation Codes

Operation Code Structure

Operation codes are 16-bit values: 0xXXYY

- XX = Category
- YY = Specific operation

Categories

Entity Operations (0x01xx)

- 0x0100 ENTITY_CREATE
- 0x0101 ENTITY_READ
- 0x0102 ENTITY_UPDATE
- 0x0103 ENTITY_DELETE \triangle IRREVERSIBLE
- 0x0104 ENTITY_LIST
- 0x0105 ENTITY_SEARCH

Google Workspace (0x03xx)

- 0x0300 GOOGLE_CONTACTS_SEARCH
- 0x0301 GOOGLE_CONTACTS_GET
- 0x0302 GOOGLE_CONTACTS_CREATE
- 0x0310 GOOGLE_CALENDAR_LIST
- 0x0320 GOOGLE_DRIVE_SEARCH
- 0x0330 GOOGLE_GMAIL_SEARCH
- 0x0332 GOOGLE_GMAIL_SEND \triangle IRREVERSIBLE

Communication (0x034x-0x036x)

- 0x0340 EMAIL_SEND \triangle IRREVERSIBLE
- 0x0350 SMS_SEND \triangle IRREVERSIBLE
- 0x0360 WEBHOOK_CALL \triangle IRREVERSIBLE

HTTP (0x04xx)

- 0x0400 HTTP_GET
- 0x0401 HTTP_POST
- 0x0402 HTTP_PUT
- 0x0403 HTTP_DELETE

Vector/RAG (0x07xx)

- 0x0700 QDRANT_SEARCH
- 0x0701 QDRANT_INDEX
- 0x0702 EMBEDDING_GENERATE

AI/LLM (0x08xx)

- 0x0800 LLM_COMPLETE
- 0x0801 LLM_CLASSIFY
- 0x0802 LLM_STRUCTURED
- 0x0810 LOCAL_MODEL_RUN

Storage (0x09xx)

- 0x0900 STORAGE_GET
- 0x0901 STORAGE_SET
- 0x0910 EVENT_LOG_APPEND
- 0x0911 EVENT_LOG_QUERY

Part 7: Command-Line Interface

pxyz compile

Compile XML workflow to binary:

```
pxyz compile --input workflow.xml --output graph.bin --audit --strict
```

Options:

- `--input FILE` - Input workflow.xml
- `--output FILE` - Output graph.bin
- `--audit` - Generate audit.json with metadata
- `--strict` - Treat warnings as errors

pxyz inspect

Inspect compiled binary:

```
pxyz inspect --input graph.bin --format mermaid
```

Formats:

- `text` - Human-readable summary
- `json` - Detailed metadata
- `mermaid` - Visual flowchart

pxyz check

Validate workflow without compiling:

```
pxyz check workflow.xml --strict
```

pxyz init

Create new project:

```
pxyz init --name my-workflow
```

Part 8: Design Philosophy

Why These Constraints?

Business Logic as Data

Benefit: Complete static analysis before execution

- Graph is fully traversable without execution
- All possible paths identifiable at compile time
- Behavior is deterministic and repeatable

Bounded Execution

Benefit: Guaranteed termination

- No unbounded loops or recursion
- Execution time is predictable
- Resource exhaustion is impossible

Explicit I/O

Benefit: Full auditability

- Every external interaction is declared
- Side effects are traceable
- Accidental operations are caught

Compile-Time Safety

Benefit: Problems caught early

- Dangerous patterns prevented at compilation
- No runtime surprises in production
- Constraints enforced by architecture

Human Gates

Benefit: Critical actions are controlled

- Irreversible operations require approval
- AI cannot independently perform high-risk actions
- Humans remain in control

Auditable Runtime

Benefit: Complete transparency

- ~600 lines of WebAssembly can be reviewed in hours
- Formal methods verification possible
- No hidden behaviors

Part 9: When to Use PXYZ

Ideal Use Cases ✓

- AI agents executing workflows
- Multi-step approval processes
- Data validation pipelines
- Third-party integrations
- Systems requiring audit trails
- Workflows with irreversible actions
- Any system where safety is critical

Not Ideal ✗

- Real-time high-frequency operations
- Complex algorithms (use external services)
- Stateful long-running processes
- Systems with extreme performance requirements

Part 10: Getting Started

1. Define Your Workflow

Create `workflow.xml` with nodes and edges

2. Compile

```
pxyz compile --input workflow.xml --output graph.bin --audit
```

3. Inspect

```
pxyz inspect --input graph.bin --format mermaid
```

4. Review

Check `audit.json` for validation results

5. Test

Load `graph.bin` in runtime with test data

6. Deploy

Embed `graph.bin` in production with host implementation

Appendix: Quick Reference

Node Kinds (0-6)

0=Transform, 1=External, 2=Render, 3=Signal, 4=Auth, 5=Terminal, 6=Error

Node Flags

- Bit 0: ASYNC
- Bit 1: REQUIRES_AUTH
- Bit 2: HAS_SIDE_EFFECTS
- Bit 3: IRREVERSIBLE
- Bit 4: REQUIRES_HUMAN
- Bit 5: CACHEABLE

Edge Flags

- Bit 0: PARALLEL
- Bit 1: FALBACK
- Bit 2: ERROR_EDGE

Safety Limits

- MAX_VISITED_NODES: 1000
- MAX_PREDICATE_STEPS: 256
- MAX_STACK_DEPTH: 16
- MAX_CALL_DEPTH: 4
- MAX_PREDICATE_BYTECODE: 256 bytes

Magic Number

0x504E5958 (ASCII: "PXYZ")

PXYZ: Safely Grant AI Agents Access to Production Systems

**